S. B. JAIN INSTITUTE OF TECHNOLOGY, MANAGEMENT & RESEARCH, NAGPUR

Practical No. 1(a)

Aim: Implement the concept of Linear Regression Models in Machine Learning with Kaggle dataset.

Name of Student : Shrutika Pradeep Bagdi

Roll No. : CS22130

Semester/Year : 6th / 3rd

Academic Session : 2024-2025

Date of Performance :

Date of Submission :

AIM: Implement the concept of Linear Regression Models in Machine Learning.

OBJECTIVE/EXPECTED LEARNING OUTCOME:

The objectives and expected learning outcome of this practical are:

- Develop a deeper understanding of the linear regression model and its limitations;
- Know how to diagnose and apply corrections to some problems with the generalized linear model found in real data; discussed;
- Use and understand generalizations of the linear model to binary and count data;
- Develop a greater familiarity with a range of techniques and methods through a diverse set of theoretical and applied readings;
- Know where to go to learn more about the techniques in this class and those called for that were not covered in this class.

THEORY:

Linear Regression:

It is the basic and commonly used type for predictive analysis. It is a statistical approach to modeling the relationship between a dependent variable and a given set of independent variables. These are of two types:

1) Simple linear Regression

2) Multiple Linear Regressions

Linear Regression is a machine learning algorithm based on supervised learning. It performs a regression task. Regression models a target prediction value based on independent variables. It is mostly used for finding out the relationship between variables and forecasting. Different regression models differ based on – the kind of relationship between dependent and independent variables they are considering, and the number of independent variables getting used. There are many names for a regression's dependent variable. It may be called an outcome variable,

criterion variable, endogenous variable, or regressand. The independent variables can be called exogenous variables, predictor variables, or regressors.

Linear regression is used in many different fields, including finance, economics, and psychology, to understand and predict the behavior of a particular variable. For example, in finance, linear regression might be used to understand the relationship between a company's stock price and its earnings, or to predict the future value of a currency based on its past performance.

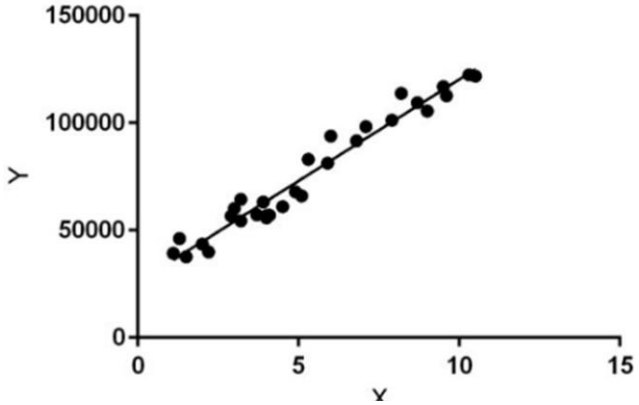

Linear regression performs the task to predict a dependent variable value (y) based on a given independent variable (x)). Hence, the name is Linear Regression. In the figure above, X (input) is the work experience and Y (output) is the salary of a person. The regression line is the best fit line for our model.

Sample Code(For reference only):

```
import numpy as np
import matplotlib as mpl
from mpl_toolkits.mplot3d import Axes3D
import matplotlib.pyplot as plt

def generate_dataset(n):
    x = []
    y = []
    random_x1 = np.random.rand()
    random_x2 = np.random.rand()
    for i in range(n):
        x1 = i
        x2 = i/2 + np.random.rand()*n
        x.append([1, x1, x2])
        y.append(random_x1 * x1 + random_x2 * x2 + 1)
```

```
    return np.array(x), np.array(y)

x, y = generate_dataset(200)

mpl.rcParams['legend.fontsize'] = 12

fig = plt.figure()
ax = fig.add_subplot(projection ='3d')

ax.scatter(x[:, 1], x[:, 2], y, label ='y', s = 5)
ax.legend()
ax.view_init(45, 0)

plt.show()
```

OUTPUT (SCREENSHOT):

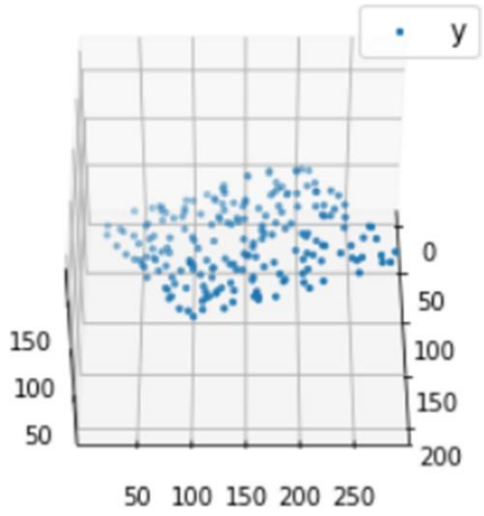

PROGRAM CODE:

OUTPUT (SCREENSHOT):

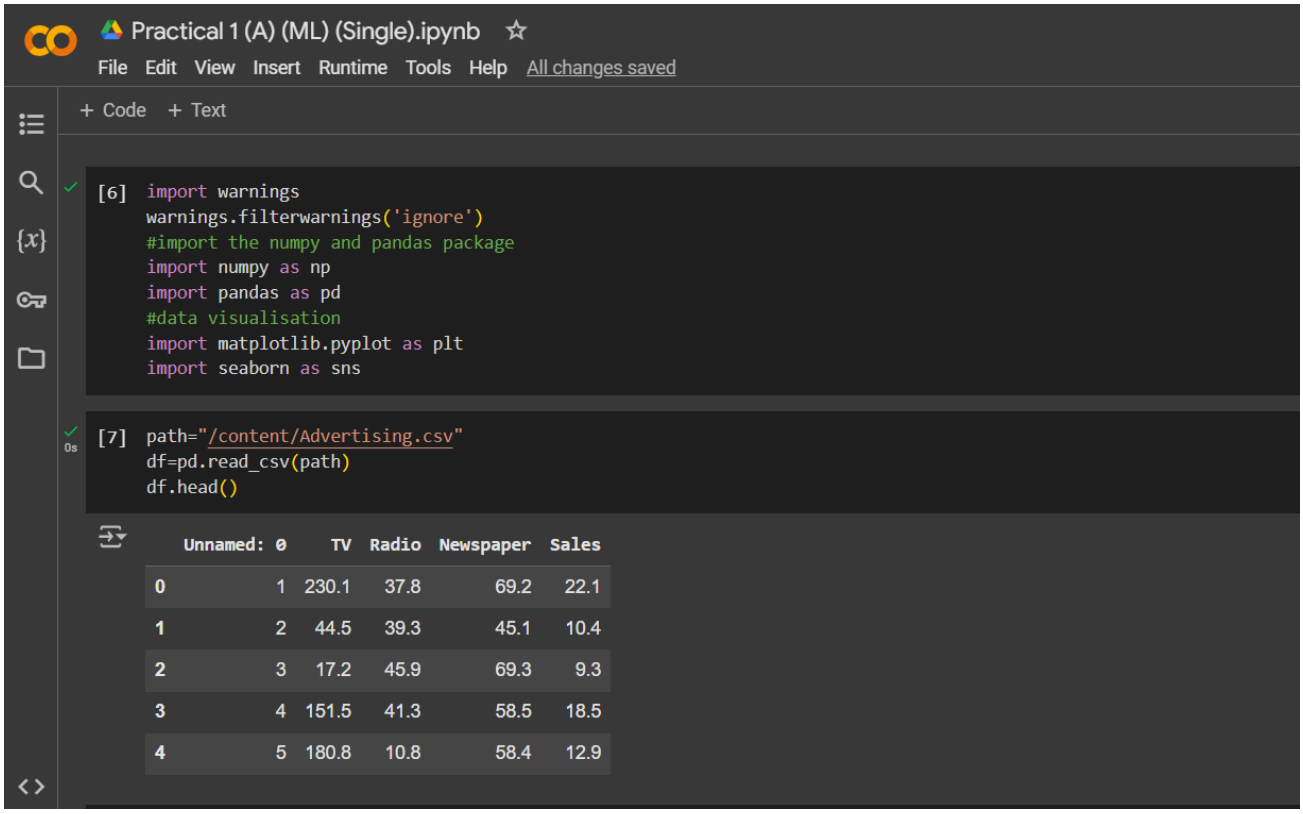

Practical 1 (A) (ML) (Single).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[8] df
```

	Unnamed: 0	TV	Radio	Newspaper	Sales
0	1	230.1	37.8	69.2	22.1
1	2	44.5	39.3	45.1	10.4
2	3	17.2	45.9	69.3	9.3
3	4	151.5	41.3	58.5	18.5
4	5	180.8	10.8	58.4	12.9
...	...	...	...	...	...
195	196	38.2	3.7	13.8	7.6
196	197	94.2	4.9	8.1	9.7
197	198	177.0	9.3	6.4	12.8
198	199	283.6	42.0	66.2	25.5
199	200	232.1	8.6	8.7	13.4

200 rows × 5 columns

```
df.shape
```

```
(200, 5)
```

Practical 1 (A) (ML) (Single).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[10] df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 200 entries, 0 to 199
Data columns (total 5 columns):
 #   Column      Non-Null Count  Dtype
---  ------      --------------  -----
 0   Unnamed: 0  200 non-null    int64
 1   TV          200 non-null    float64
 2   Radio       200 non-null    float64
 3   Newspaper   200 non-null    float64
 4   Sales       200 non-null    float64
dtypes: float64(4), int64(1)
memory usage: 7.9 KB
```

```
df.describe()
```

	Unnamed: 0	TV	Radio	Newspaper	Sales
count	200.000000	200.000000	200.000000	200.000000	200.000000
mean	100.500000	147.042500	23.264000	30.554000	14.022500
std	57.879185	85.854236	14.846809	21.778621	5.217457
min	1.000000	0.700000	0.000000	0.300000	1.600000
25%	50.750000	74.375000	9.975000	12.750000	10.375000
50%	100.500000	149.750000	22.900000	25.750000	12.900000
75%	150.250000	218.825000	36.525000	45.100000	17.400000
max	200.000000	296.400000	49.600000	114.000000	27.000000

Practical 1 (A) (ML) (Single).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[12] #let's see how sales are related with other variable using  sctter plot.
     #sns.piarplot(df,x_vars=['TV','Newspaper','radio'],y_vars='sales',height=4,aspect=height,aspect)
     sns.pairplot(df,x_vars=['TV'],y_vars='Sales',height=4,aspect=1,kind='scatter')
     plt.show()
```

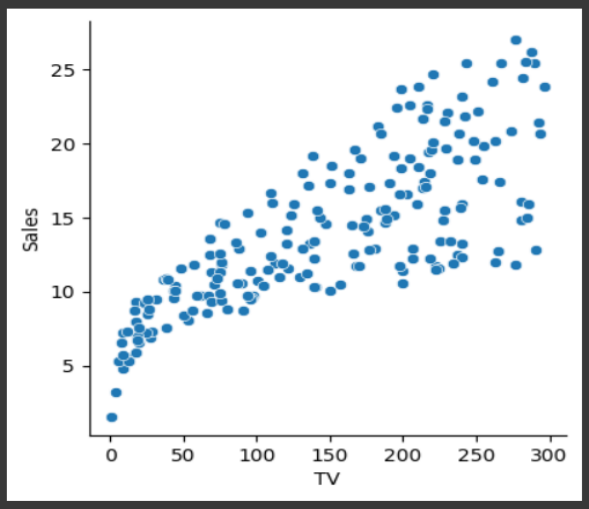


```
x=df['TV']
y=df['Sales']
```

Practical 1 (A) (ML) (Single).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[14] from sklearn.model_selection import train_test_split
     x_train,x_test,y_train,y_test=train_test_split(x,y,train_size = 0.7,test_size=0.3,random_state=100)
```

```
x_train
```

	TV
74	213.4
3	151.5
185	205.0
26	142.9
90	134.3
...	...
87	110.7
103	187.9
67	139.3
24	62.3
8	8.6

140 rows × 1 columns

dtype: float64

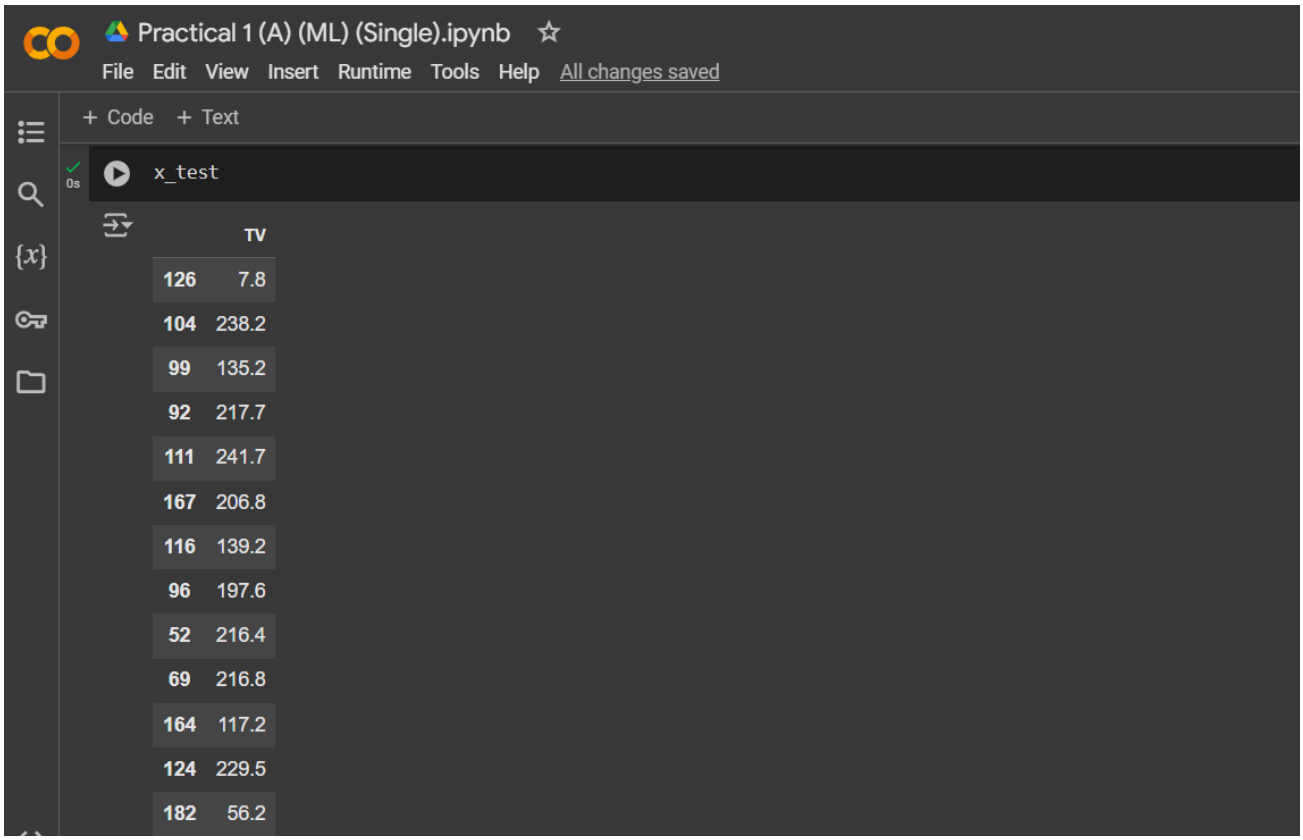
Practical 1 (A) (ML) (Single).ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
x_test
TV
126 7.8
104 238.2
99 135.2
92 217.7
111 241.7
167 206.8
116 139.2
96 197.6
52 216.4
69 216.8
164 117.2
124 229.5
182 56.2

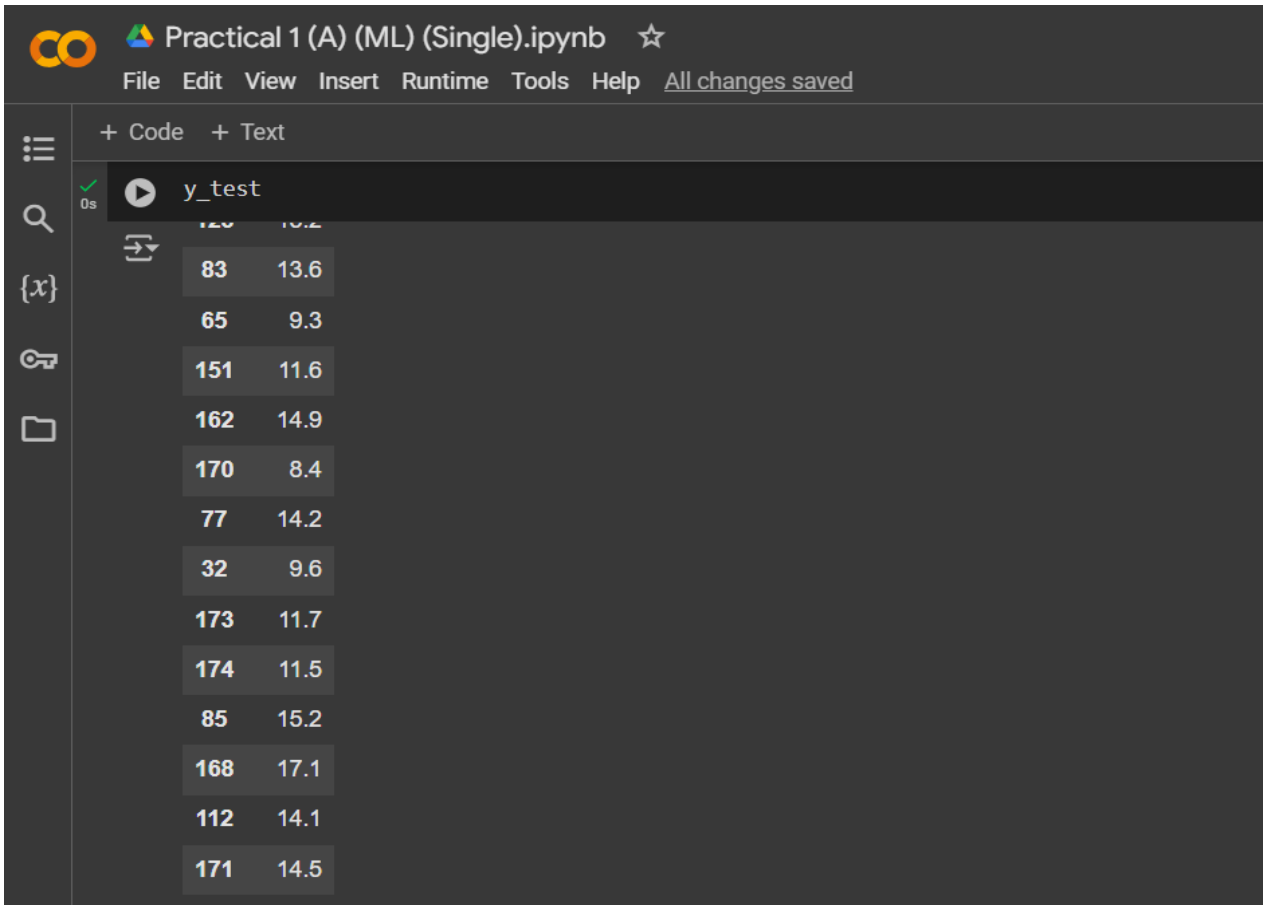
Practical 1 (A) (ML) (Single).ipynb
File Edit View Insert Runtime Tools Help All changes saved
+ Code + Text
y_test
83 13.6
65 9.3
151 11.6
162 14.9
170 8.4
77 14.2
32 9.6
173 11.7
174 11.5
85 15.2
168 17.1
112 14.1
171 14.5

Practical 1 (A) (ML) (Single).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[18] y_train
```

	Sales
74	17.0
3	18.5
185	22.6
26	15.0
90	11.2
...	...
87	16.0
103	14.7
67	13.4
24	9.7
8	4.8

140 rows × 1 columns

dtype: float64

```
[19] import statsmodels.api as sm
```

Practical 1 (A) (ML) (Single).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[20] x_train_sm=sm.add_constant(x_train)

     lr=sm.OLS(y_train,x_train_sm).fit()
```

```
[21] lr.params
```

	0
const	6.989666
TV	0.046497

dtype: float64

Practical 1 (A) (ML) (Single).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[22] print(lr.summary())
```

```
                            OLS Regression Results
==============================================================================
Dep. Variable:                  Sales   R-squared:                       0.613
Model:                            OLS   Adj. R-squared:                  0.611
Method:                 Least Squares   F-statistic:                     219.0
Date:                Sun, 12 Jan 2025   Prob (F-statistic):           2.84e-30
Time:                        18:14:57   Log-Likelihood:                -370.62
No. Observations:                 140   AIC:                             745.2
Df Residuals:                     138   BIC:                             751.1
Df Model:                           1
Covariance Type:            nonrobust
==============================================================================
                 coef    std err          t      P>|t|      [0.025      0.975]
------------------------------------------------------------------------------
const          6.9897      0.548     12.762      0.000       5.907       8.073
TV             0.0465      0.003     14.798      0.000       0.040       0.053
==============================================================================
Omnibus:                        0.995   Durbin-Watson:                   1.983
Prob(Omnibus):                  0.608   Jarque-Bera (JB):                0.970
Skew:                          -0.008   Prob(JB):                        0.616
Kurtosis:                       2.593   Cond. No.                         328.
==============================================================================

Notes:
[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.
```

Practical 1 (A) (ML) (Single).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
plt.scatter(x_train,y_train,c='r')
plt.plot(x_train,6.9897+0.0465*x_train,'y')
plt.show()
```

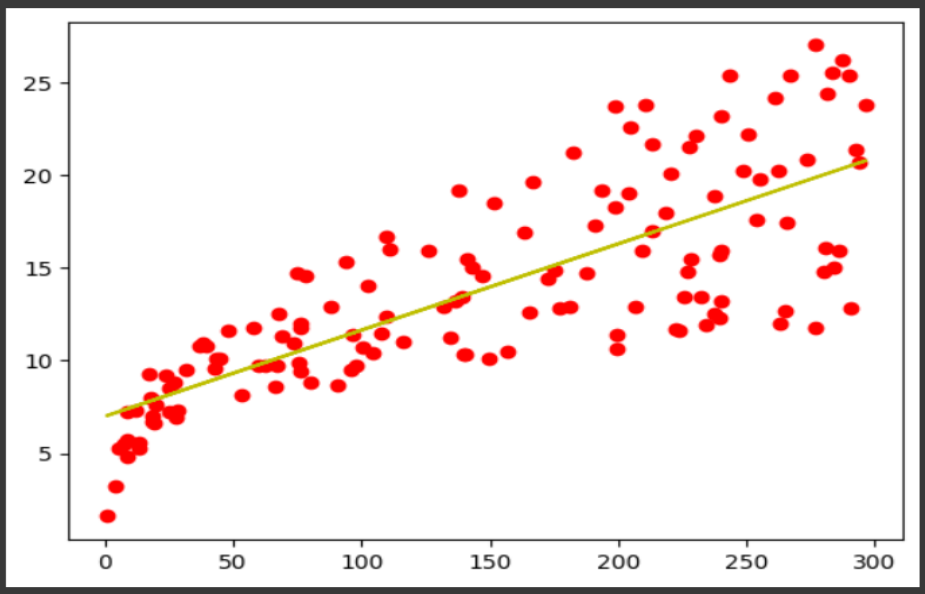

Practical 1 (A) (DMW) (Multiple).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
sns.pairplot(df,x_vars=['TV'],y_vars='Sales',height=4,aspect=1,kind='scatter')
plt.show()
```

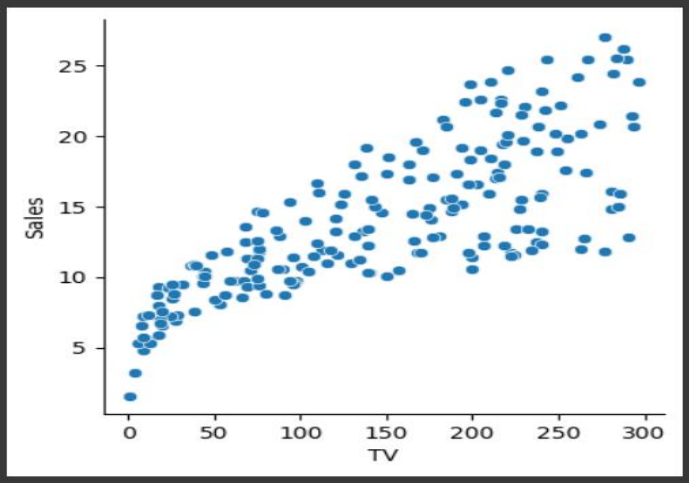


```
[8] x=df[['TV','Radio','Newspaper']]
    y=df['Sales']
```

Practical 1 (A) (DMW) (Multiple).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[10] x.head()
```

	TV	Radio	Newspaper
0	230.1	37.8	69.2
1	44.5	39.3	45.1
2	17.2	45.9	69.3
3	151.5	41.3	58.5
4	180.8	10.8	58.4

```
[11] from sklearn.model_selection import train_test_split
     x_train,x_test,y_train,y_test=train_test_split(x,y,train_size = 0.7,test_size=0.3,random_state=100)
```

```
x_train
```

	TV	Radio	Newspaper
74	213.4	24.6	13.1
3	151.5	41.3	58.5
185	205.0	45.1	19.6
26	142.9	29.3	12.6
90	134.3	4.9	9.3
...	...	...	...

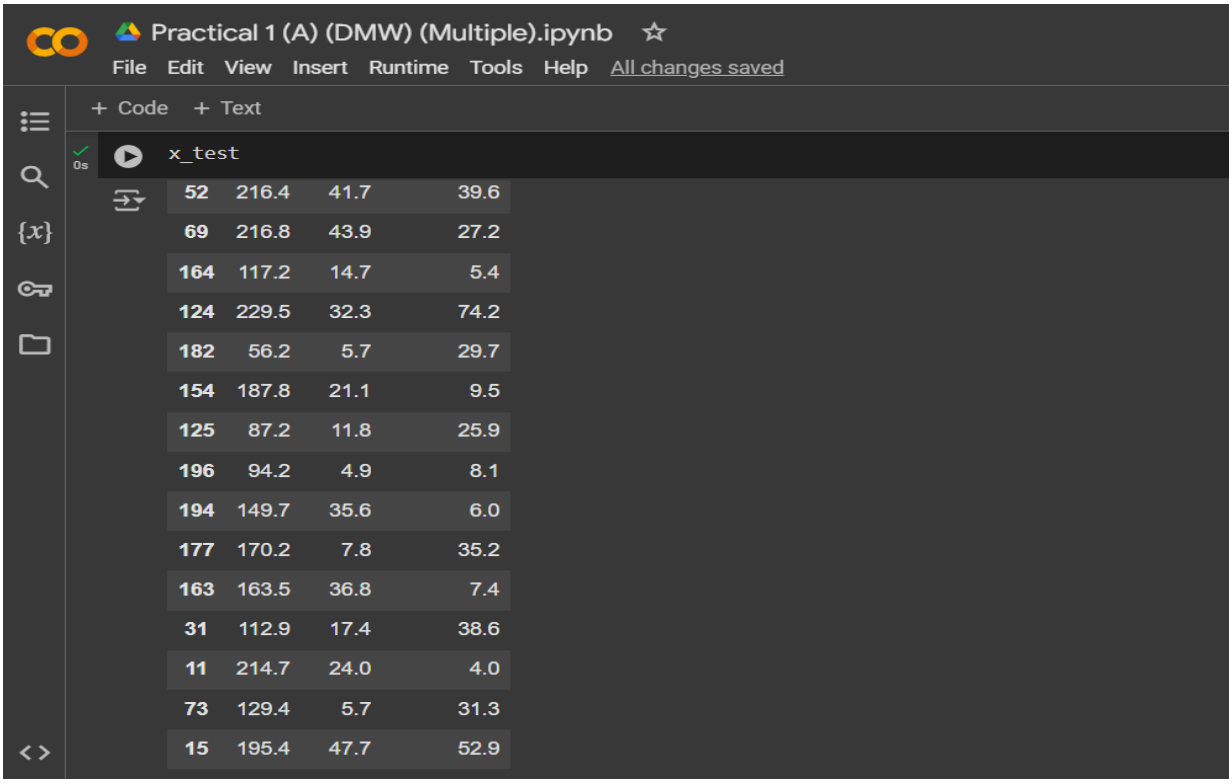

52	216.4	41.7	39.6
69	216.8	43.9	27.2
164	117.2	14.7	5.4
124	229.5	32.3	74.2
182	56.2	5.7	29.7
154	187.8	21.1	9.5
125	87.2	11.8	25.9
196	94.2	4.9	8.1
194	149.7	35.6	6.0
177	170.2	7.8	35.2
163	163.5	36.8	7.4
31	112.9	17.4	38.6
11	214.7	24.0	4.0
73	129.4	5.7	31.3
15	195.4	47.7	52.9

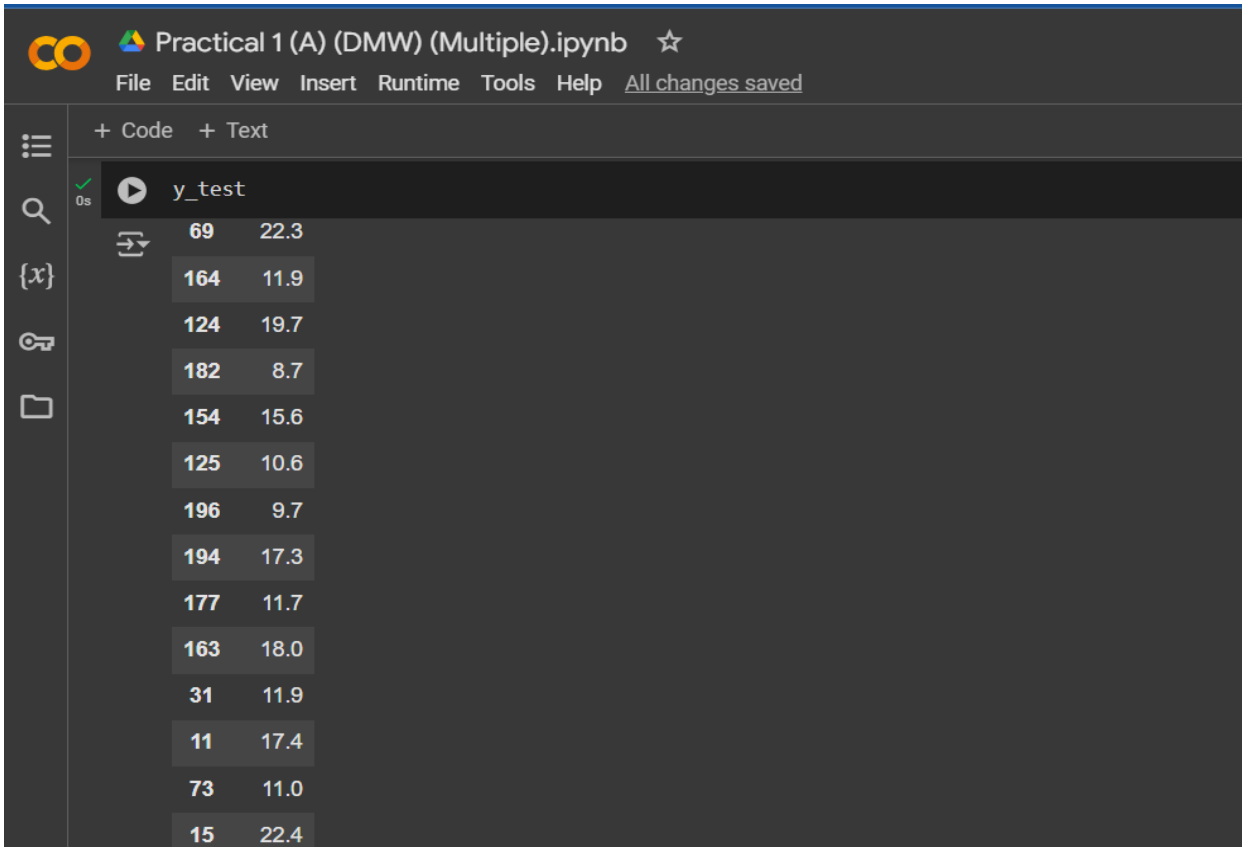

69	22.3
164	11.9
124	19.7
182	8.7
154	15.6
125	10.6
196	9.7
194	17.3
177	11.7
163	18.0
31	11.9
11	17.4
73	11.0
15	22.4

Practical 1 (A) (DMW) (Multiple).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[15] y_train
```

	Sales
74	17.0
3	18.5
185	22.6
26	15.0
90	11.2
...	...
87	16.0
103	14.7
67	13.4
24	9.7
8	4.8

140 rows × 1 columns

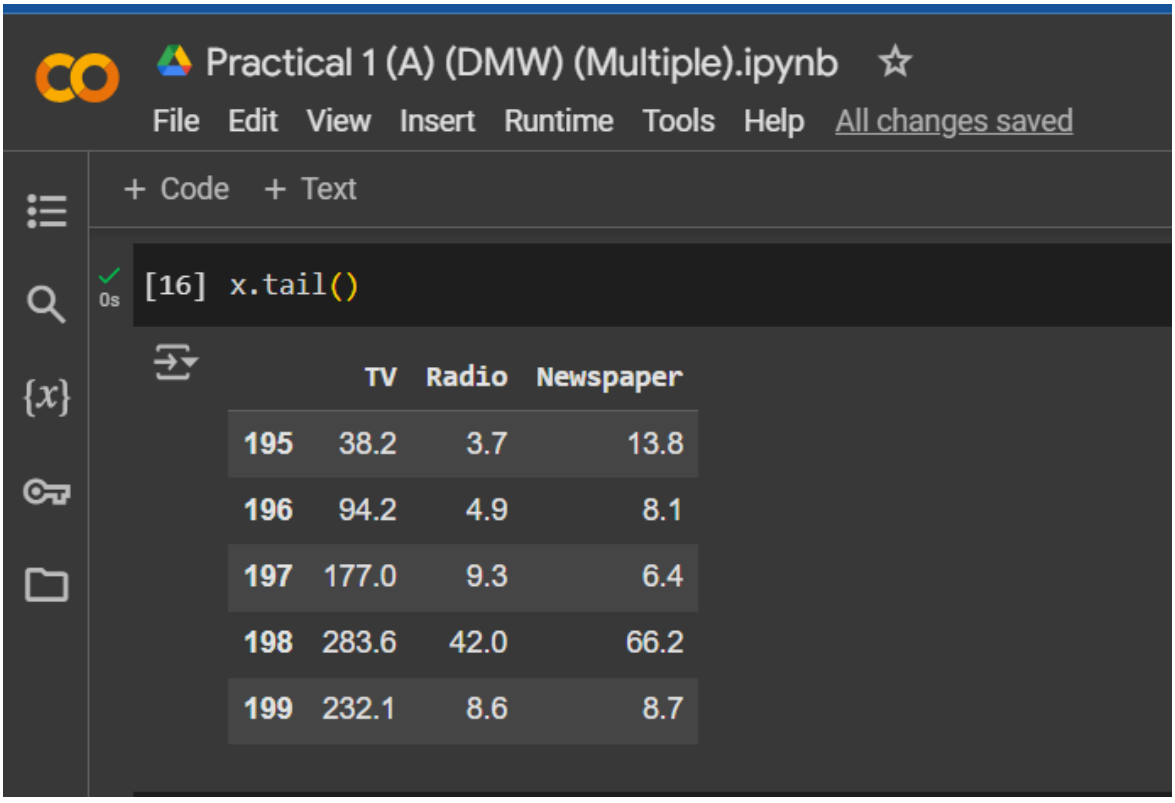

	TV	Radio	Newspaper
195	38.2	3.7	13.8
196	94.2	4.9	8.1
197	177.0	9.3	6.4
198	283.6	42.0	66.2
199	232.1	8.6	8.7

Practical 1 (A) (DMW) (Multiple).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
import seaborn as sns
sns.pairplot(df,x_vars=['TV','Newspaper','Radio'],y_vars='Sales',height=4,aspect=1,kind='scatter')
```

```
<seaborn.axisgrid.PairGrid at 0x7a89a47e13f0>
```

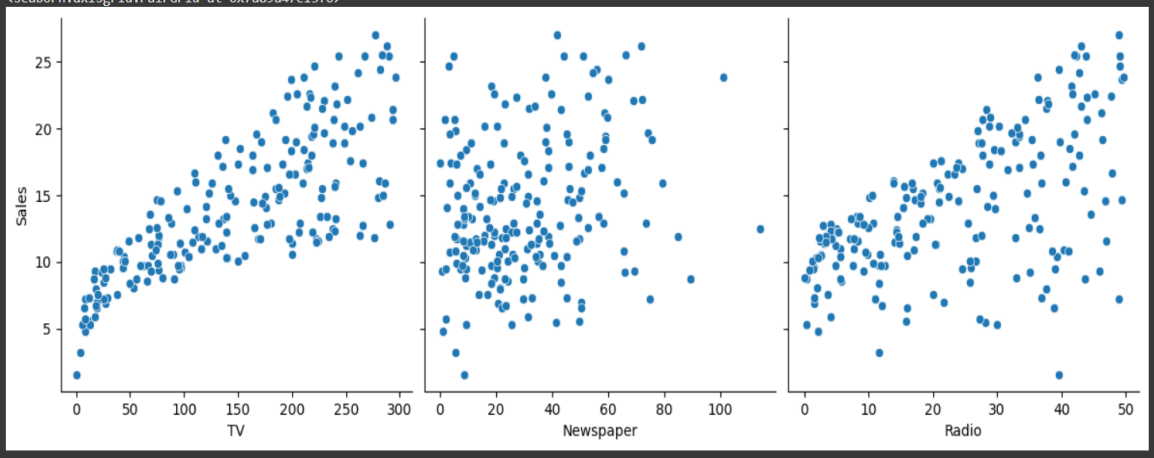


```
[18] from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test=train_test_split(x,y,train_size = 0.7,test_size=0.3,random_state=100)
```

Practical 1 (A) (DMW) (Multiple).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[19] x_train.head()
```

	TV	Radio	Newspaper
74	213.4	24.6	13.1
3	151.5	41.3	58.5
185	205.0	45.1	19.6
26	142.9	29.3	12.6
90	134.3	4.9	9.3

```
y_train.head()
```

	Sales
74	17.0
3	18.5
185	22.6
26	15.0
90	11.2

dtype: float64

Practical 1 (A) (DMW) (Multiple).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
[21] from sklearn.linear_model import LinearRegression
     regression=LinearRegression()
     regression.fit(x_train,y_train)
     reg_prediction=regression.predict(x_test)
```

```
[22] reg_prediction
```

```
array([10.62160072, 20.00625302, 16.91850882, 19.17040746, 20.94974131,
       13.12284284, 11.80740696, 12.32019766, 20.57806782, 20.95662688,
       10.79096475, 19.54868702,  6.42403866, 15.23133391,  8.97226257,
        7.89897862, 16.23599497, 12.02636477, 17.09702178, 11.26080277,
       16.97826292,  9.75655721, 20.82389762, 17.20916742, 15.13816239,
       21.97290698, 19.20181841, 10.07501899, 19.39017185, 14.8673761 ,
       14.36798893,  7.55604543,  9.96742165, 14.76342565,  7.20995576,
       13.60003295,  7.49088656, 11.70865932, 13.46091883, 15.2229793 ,
       17.18088277, 13.56738329, 14.30942267, 13.72909849, 11.88559349,
        8.77039705, 12.1244102 , 19.20252289,  9.08376601,  5.15367352,
       16.22852749, 18.14111213, 12.94835466, 16.86274503, 17.86462435,
       12.33930625,  4.3575739 , 11.25904494, 16.11560622, 13.56602169])
```

```
import seaborn as sns
sns.displot(reg_prediction-y_test)
```

Practical 1 (A) (DMW) (Multiple).ipynb

File Edit View Insert Runtime Tools Help All changes saved

+ Code + Text

```
sns.displot(reg_prediction-y_test,kind='kde')
```

```
<seaborn.axisgrid.FacetGrid at 0x7a899fd410f0>
```

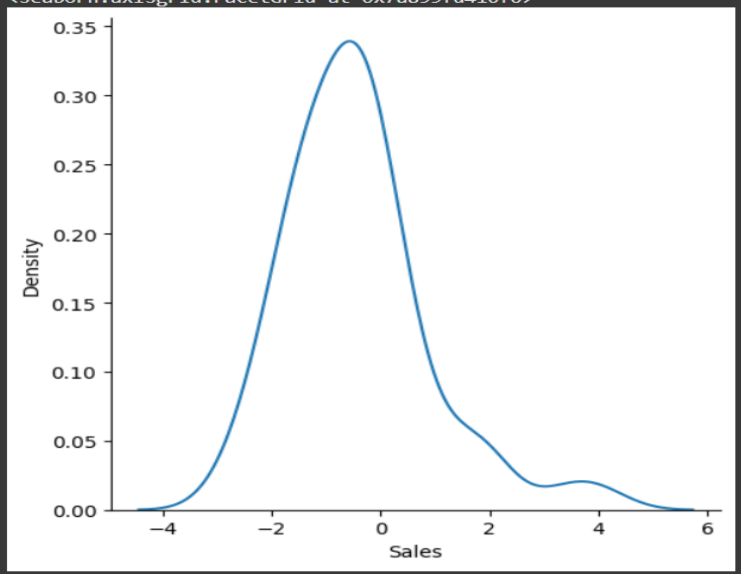

CONCLUSION:

_ _

_ _

_ _

_ _

_ _

DISCUSSION AND VIVA VOCE:

- What is linear regression in machine learning example?
- What algorithm is used in linear regression?
- What is the application of linear regression?
- What are some real life examples of linear regression?
- What is linear regression and its types?

REFERENCE:

- https://www.google.com/search?q=linear+regression+machine+learning&oq=Linear+Regression&aqs=chrome.1.0i433i512j0i512j69i59j0i131i433i512l2j69i60l3.4688j0j7&sourceid=chrome&ie=UTF-8
- https://pubs.acs.org/doi/10.1021/acsomega.2c00362
- https://colab.research.google.com/drive/1gbyjW-RgdgVkUNyRkHAWmjMr7RUCUrYX